

Documentation for the "Best Time to Buy and Sell Stock with Cooldown"

Problem

The problem requires finding the maximum profit achievable from buying and selling stocks over multiple days, subject to the constraint that after selling a stock, you must wait one day (cooldown) before buying again. The goal is to determine an optimal strategy to maximize profits while adhering to these rules.

Problem Breakdown and Constraints

The input is an array of prices where each element represents the stock price on a specific day. You can perform as many buy-and-sell operations as you like, but you must sell any stock before buying a new one. Additionally, you cannot buy a stock on the day immediately following a sale. The problem assumes that the length of prices lies between 1 and 5000, and stock prices range from 0 to 1000. These constraints ensure that the problem is computationally feasible using dynamic programming techniques.

Key Insights and Approach

The core insight to solving the problem involves tracking three distinct states for each day:

1. **Hold:** The maximum profit achievable if holding a stock at the end of the current day.
2. **sold:** The maximum profit achievable if selling a stock on the current day.
3. **rest:** The maximum profit achievable if neither buying nor selling on the current day (cooldown or idle state).

By defining these states, the solution can compute optimal values for each day based on the previous day's states. This method ensures that the problem is broken down into smaller, manageable subproblems, making it ideal for a dynamic programming approach.

Transition Relations

The transitions between states capture the dynamics of buying, selling, or resting. For example, the hold state on any day depends on whether the stock was already being held or was bought from the rest state on that day. Similarly, the sold state depends on selling a stock held from the previous day, and the rest state depends on either staying in the rest state or entering it after selling. These transitions provide a systematic way to update states daily and maximize profits.

Base Cases and Final Results

For day 0 (the first day), the initial states are straightforward:

- hold is initialized to $-\text{prices}[0]$ since buying the stock incurs a cost.
- sold and rest are initialized to 0 because no profit can be made without prior transactions.

The maximum profit at the end of the last day is determined by taking the maximum of the sold and rest states, as holding a stock on the last day cannot maximize profit.

Computational Complexity

The solution iterates through the prices array once, making it linear in time complexity ($O(n)$), where n is the number of days. Space complexity can be optimized to $O(1)$ by using only variables to track the states of the previous day instead of maintaining arrays for all days. This efficiency ensures that the solution works well within the given constraints.

Practical Applications

The approach used in this problem is a classic example of state-based decision-making. It is applicable in real-world scenarios where certain constraints, such as cooldown periods or restrictions on consecutive actions, must be considered. This methodology is also relevant in financial modeling, portfolio optimization, and game theory, where optimal strategies are derived based on dynamic conditions.

Challenges and Edge Cases

While the solution works effectively for standard scenarios, edge cases must be handled appropriately. For instance, if the array length is 1, no transactions can occur, and the profit is zero. Also, arrays with constant or non-increasing prices should return zero profit, as no beneficial transactions are possible. The solution accounts for these cases to ensure correctness across all possible inputs.