

Documentation for Sliding Window Median

1. Problem Statement

You are given an array of integers, `nums`, and an integer `k`. There is a sliding window of size `k` moving from left to right of the array. At each step, the window moves right by one position.

You must return an array of medians where `medians[i]` is the median of the `i`-th sliding window.

- If `k` is odd, the median is the middle element.
- If `k` is even, the median is the average of the two middle elements.

Example:

Input: `nums = [1,3,-1,-3,5,3,6,7]`, `k = 3`

Output: `[1.00000, -1.00000, -1.00000, 3.00000, 5.00000, 6.00000]`

2. Intuition

To maintain the median efficiently as the window slides:

- We need a way to insert and remove elements in \log time.
- We also need to quickly access the middle elements.

This motivates the use of two heaps:

- A max heap (small) for the smaller half of the window.
- A min heap (large) for the larger half.

3. Key Observations

- The max heap stores the negative values because Python only supports min heaps.
- For a window of size `k`:
 - If `k` is odd, the median is the top of the small.
 - If `k` is even, it's the average of the tops of small and large.

- Elements exiting the window need to be delayed deleted, not immediately removed from the heap.
- We use a hash map (delayed) to keep track of elements that should be removed when they reach the top of a heap.

4. Approach

Steps:

- Use two heaps: small (max heap) and large (min heap).
- Maintain heap balance such that:
 - Small has the same size as large or one more.
- Insert new elements into one of the heaps.
- Remove elements that have exited the window using delayed deletion.
- Prune the top of heaps whenever invalid elements are found.
- After each valid window, compute the median.

5. Edge Cases

- nums with negative and positive values.
- Repeated values in nums.
- When $k == 1$, the median is simply the element itself.
- When $k == \text{len}(\text{nums})$, only one median is returned.
- Large values (up to $2^{31}-1$) must be handled efficiently.

6. Complexity Analysis

□ Time Complexity:

- Each insertion/removal takes $O(\log k)$.
- Total operations = n (length of nums).
- So, overall time complexity: $O(n \log k)$.

📦 Space Complexity:

- Two heaps of size up to k : $O(k)$
- Delayed map for at most k items: $O(k)$
- Total space complexity: $O(k)$

7. Alternative Approaches

✖ Brute Force:

- Sort every window: $O(k \log k)$ per window.
- Total: $O(nk \log k)$ — too slow for large inputs.

✓ Balanced BST (like SortedList from bisect or SortedContainers):

- Insertion/removal in $O(\log k)$
- Easier than heap-based with delayed deletions.
- Cleaner but requires external libraries or more complex code in pure Python.

8. Test Cases

```
assert Solution().medianSlidingWindow([1,3,-1,-3,5,3,6,7], 3) == [1.0, -1.0, -1.0, 3.0, 5.0, 6.0]
assert Solution().medianSlidingWindow([1,2,3,4,2,3,1,4,2], 3) == [2.0, 3.0, 3.0, 3.0, 2.0, 3.0, 2.0]
assert Solution().medianSlidingWindow([1], 1) == [1.0]
assert Solution().medianSlidingWindow([1, 2], 2) == [1.5]
```

9. Final Thoughts

- The two-heaps approach is efficient and optimal for this type of median tracking problem.
- It's a great example of using data structures cleverly to maintain invariants (like balance and order).
- The delayed deletion strategy may seem tricky at first but is necessary due to heap limitations in Python.

